# Parallel Computing Benchmark Document — Page 1 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 2 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 3 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 4 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 5 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 6 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 7 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 8 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 9 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 10 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 11 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 12 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 13 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

# Parallel Computing Benchmark Document — Page 14 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 15 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 16 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 17 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 18 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 19 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

# Parallel Computing Benchmark Document — Page 20 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 21 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 22 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 23 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

# Parallel Computing Benchmark Document — Page 24 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 25 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 26 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 27 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 28 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 29 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

# Parallel Computing Benchmark Document — Page 30 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 31 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 32 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 33 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 34 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 35 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 36 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 37 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 38 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 39 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 40 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 41 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 42 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 43 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

# Parallel Computing Benchmark Document — Page 44 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 45 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 46 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 47 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 48 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 49 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

# Parallel Computing Benchmark Document — Page 50 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 51 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 52 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 53 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

# Parallel Computing Benchmark Document — Page 54 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 55 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 56 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 57 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 58 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 59 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

# Parallel Computing Benchmark Document — Page 60 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 61 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 62 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 63 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 64 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 65 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 66 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 67 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 68 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 69 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 70 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 71 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 72 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 73 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

# Parallel Computing Benchmark Document — Page 74 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 75 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 76 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 77 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 78 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 79 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

# Parallel Computing Benchmark Document — Page 80 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 81 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 82 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 83 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

# Parallel Computing Benchmark Document — Page 84 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 85 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 86 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 87 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 88 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 89 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

# Parallel Computing Benchmark Document — Page 90 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 91 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

# Parallel Computing Benchmark Document — Page 92 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 93 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 94 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 95 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

# Parallel Computing Benchmark Document — Page 96 of 100

## Amdahl's Law and Theoretical Speedup Bounds (Section 1)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 2)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 3)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 4)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 97 of 100

## Process Scheduling and Concurrency in Operating Systems (Section 1)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs.

## Data Partitioning and Load Balancing (Section 2)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest.

## Inter-Process Communication and Overhead (Section 3)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput.

## Parallel Computing and High Performance Systems (Section 4)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes.

# Parallel Computing Benchmark Document — Page 98 of 100

## Data Partitioning and Load Balancing (Section 1)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Inter-Process Communication and Overhead (Section 2)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 3)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 4)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 99 of 100

## Inter-Process Communication and Overhead (Section 1)

Inter-process communication (IPC) mechanisms allow processes to manage shared data. However, excessive serialization (pickling) and IPC bandwidth can become a significant bottleneck. In our PDF extractor design, each worker independently reads its assigned page range from disk, eliminating redundant memory transfers across process boundaries and boosting throughput. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Parallel Computing and High Performance Systems (Section 2)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 3)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 4)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.

# Parallel Computing Benchmark Document — Page 100 of 100

## Parallel Computing and High Performance Systems (Section 1)

Parallel computing is a type of computation in which many calculations or processes are carried out simultaneously. Large problems can often be divided into smaller ones, which can then be solved at the same time. There are several different forms of parallel computing: bit-level, instruction-level, data, and task parallelism. In modern multi-core architectures, achieving maximum speedup requires minimizing synchronization overhead and balanced workload distribution among active worker processes. Notice: The keyword 'Amdahl' is discussed in detail within Section 1 regarding system scaling and parallel efficiency benchmarks.

## Amdahl's Law and Theoretical Speedup Bounds (Section 2)

Amdahl's law is a formula which gives the theoretical speedup in latency of the execution of a task at fixed workload that can be expected of a system whose resources are improved. Named after computer scientist Gene Amdahl, it states that if a program has a serial portion s, the maximum achievable speedup using p processors is bounded by 1 / (s + (1 - s) / p). Even with an infinite number of processors, the speedup can never exceed 1 / s. Notice: The keyword 'Amdahl' is discussed in detail within Section 2 regarding system scaling and parallel efficiency benchmarks.

## Process Scheduling and Concurrency in Operating Systems (Section 3)

Concurrency is the execution of multiple instruction sequences at the same time. In modern operating systems, multiprocessing allows true parallel execution on multiple CPU cores, avoiding the limitations of the Global Interpreter Lock (GIL) found in CPython. When designing parallel algorithms, worker allocation and process creation cost must be carefully evaluated to avoid performance degradation on small computational jobs. Notice: The keyword 'Amdahl' is discussed in detail within Section 3 regarding system scaling and parallel efficiency benchmarks.

## Data Partitioning and Load Balancing (Section 4)

In distributed and shared-memory parallel computing, load balancing refers to the practice of distributing work evenly across available computing resources. A parallel efficiency metric measures the fraction of time for which processors are being used constructively. Imbalanced partitions cause worker starvation where fast workers idle while waiting for the slowest. Notice: The keyword 'Amdahl' is discussed in detail within Section 4 regarding system scaling and parallel efficiency benchmarks.